

# Notes on Tidyverse

*Shashank Sule*

*9/15/2019*

Tidyverse is the collection of tools which help you visualize and explore data. The example we will use is the gapminder data set.

```
require(gapminder)
```

```
## Loading required package: gapminder
```

```
require(dplyr)
```

```
## Loading required package: dplyr
```

```
## Warning: package 'dplyr' was built under R version 3.4.4
```

```
##
```

```
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
##      filter, lag
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
##      intersect, setdiff, setequal, union
```

```
gapminder
```

```
## # A tibble: 1,704 x 6
```

	country	continent	year	lifeExp	pop	gdpPercap
	<fct>	<fct>	<int>	<dbl>	<int>	<dbl>
## 1	Afghanistan	Asia	1952	28.8	8425333	779
## 2	Afghanistan	Asia	1957	30.3	9240934	821
## 3	Afghanistan	Asia	1962	32.0	10267083	853
## 4	Afghanistan	Asia	1967	34.0	11537966	836
## 5	Afghanistan	Asia	1972	36.1	13079460	740
## 6	Afghanistan	Asia	1977	38.4	14880372	786
## 7	Afghanistan	Asia	1982	39.9	12881816	978
## 8	Afghanistan	Asia	1987	40.8	13867957	852
## 9	Afghanistan	Asia	1992	41.7	16317921	649
## 10	Afghanistan	Asia	1997	41.8	22227415	635

```
## # ... with 1,694 more rows
```

Note that gapminder is a tibble, a bit different from data.frame.

# Verbs from dplyr

Verbs are the functions within dplyr. Every time you use a verb, you use a pipe which composes verbs.

## 1. filter

`filter` literally filters your data frame according to the criteria you give. Suppose you want to filter the dataset for 2007.

```
gapminder %>%  
  filter(year == 2007) #Indexing the array by boolean variables
```

```
## Warning: package 'bindrcpp' was built under R version 3.4.4
```

```
## # A tibble: 142 x 6
```

	country	continent	year	lifeExp	pop	gdpPercap
	<fct>	<fct>	<int>	<dbl>	<int>	<dbl>
## 1	Afghanistan	Asia	2007	43.8	31889923	975
## 2	Albania	Europe	2007	76.4	3600523	5937
## 3	Algeria	Africa	2007	72.3	33333216	6223
## 4	Angola	Africa	2007	42.7	12420476	4797
## 5	Argentina	Americas	2007	75.3	40301927	12779
## 6	Australia	Oceania	2007	81.2	20434176	34435
## 7	Austria	Europe	2007	79.8	8199783	36126
## 8	Bahrain	Asia	2007	75.6	708573	29796
## 9	Bangladesh	Asia	2007	64.1	150448339	1391
## 10	Belgium	Europe	2007	79.4	10392226	33693
##	... with 132 more rows					

```
#You can also double filter
```

```
gapminder %>% filter(year == 2007, country == "United States")
```

```
## # A tibble: 1 x 6
```

	country	continent	year	lifeExp	pop	gdpPercap
	<fct>	<fct>	<int>	<dbl>	<int>	<dbl>
## 1	United States	Americas	2007	78.2	301139947	42952

## 2. arrange

`arrange` sorts the dataset according to the column you want to arrange by. By default it arranges according to ascending order. To sort in descending order you do `arrange(desc(colname))`

```
# Sort in ascending order of lifeExp  
gapminder %>% arrange(lifeExp)
```

```
## # A tibble: 1,704 x 6
```

```
##   country      continent year lifeExp      pop gdpPercap
##   <fct>        <fct>    <int>  <dbl>    <int>    <dbl>
## 1 Rwanda      Africa    1992   23.6  7290203      737
## 2 Afghanistan Asia     1952   28.8  8425333      779
## 3 Gambia       Africa    1952   30.0   284320      485
## 4 Angola       Africa    1952   30.0  4232095     3521
## 5 Sierra Leone Africa    1952   30.3  2143249      880
## 6 Afghanistan Asia     1957   30.3  9240934      821
## 7 Cambodia     Asia     1977   31.2  6978607      525
## 8 Mozambique   Africa    1952   31.3  6446316      469
## 9 Sierra Leone Africa    1957   31.6  2295678     1004
## 10 Burkina Faso Africa    1952   32.0  4469979      543
## # ... with 1,694 more rows
```

```
# Sort in descending order of lifeExp
gapminder %>% arrange(desc(lifeExp))
```

```
## # A tibble: 1,704 x 6
##   country      continent year lifeExp      pop gdpPercap
##   <fct>        <fct>    <int>  <dbl>    <int>    <dbl>
## 1 Japan      Asia     2007   82.6 127467972   31656
## 2 Hong Kong, China Asia     2007   82.2  6980412   39725
## 3 Japan      Asia     2002   82.0 127065841   28605
## 4 Iceland    Europe    2007   81.8   301931   36181
## 5 Switzerland Europe    2007   81.7   7554661   37506
## 6 Hong Kong, China Asia     2002   81.5   6762476   30209
## 7 Australia  Oceania   2007   81.2  20434176   34435
## 8 Spain      Europe    2007   80.9  40448191   28821
## 9 Sweden     Europe    2007   80.9   9031088   33860
## 10 Israel    Asia     2007   80.7   6426679   25523
## # ... with 1,694 more rows
```

### 3. mutate

`mutate` either changes an existing column in the data frame or adds a new one. Note that it is not a “static” function: it does not make any changes to the data frame you pass as an argument. Instead, it returns a new data frame that is a result of passing the original data frame through `mutate`.

```
# Use mutate to change lifeExp to be in months
gapminder %>% mutate(lifeExp = lifeExp * 12)
```

```
## # A tibble: 1,704 x 6
##   country      continent year lifeExp      pop gdpPercap
##   <fct>        <fct>    <int>  <dbl>    <int>    <dbl>
## 1 Afghanistan Asia     1952   346  8425333      779
## 2 Afghanistan Asia     1957   364  9240934      821
```

```
## 3 Afghanistan Asia      1962      384 10267083      853
## 4 Afghanistan Asia      1967      408 11537966      836
## 5 Afghanistan Asia      1972      433 13079460      740
## 6 Afghanistan Asia      1977      461 14880372      786
## 7 Afghanistan Asia      1982      478 12881816      978
## 8 Afghanistan Asia      1987      490 13867957      852
## 9 Afghanistan Asia      1992      500 16317921      649
## 10 Afghanistan Asia     1997      501 22227415      635
## # ... with 1,694 more rows
```

```
# Use mutate to create a new column called lifeExpMonths
```

```
gapminder %>% mutate(lifeExpMonths = lifeExp * 12)
```

```
## # A tibble: 1,704 x 7
##   country      continent  year lifeExp      pop gdpPercap lifeExpMonths
##   <fct>         <fct>    <int> <dbl>    <int>    <dbl>      <dbl>
## 1 Afghanistan Asia      1952   28.8  8425333    779        346
## 2 Afghanistan Asia      1957   30.3  9240934    821        364
## 3 Afghanistan Asia      1962   32.0 10267083    853        384
## 4 Afghanistan Asia      1967   34.0 11537966    836        408
## 5 Afghanistan Asia      1972   36.1 13079460    740        433
## 6 Afghanistan Asia      1977   38.4 14880372    786        461
## 7 Afghanistan Asia      1982   39.9 12881816    978        478
## 8 Afghanistan Asia      1987   40.8 13867957    852        490
## 9 Afghanistan Asia      1992   41.7 16317921    649        500
## 10 Afghanistan Asia     1997   41.8 22227415    635        501
## # ... with 1,694 more rows
```

```
# Filter, mutate, and arrange the gapminder dataset
```

```
gapminder %>%
  filter(year == 2007) %>%
  mutate(lifeExpMonths = 12 * lifeExp)%>%
  arrange(desc(lifeExpMonths))
```

```
## # A tibble: 142 x 7
##   country      continent  year lifeExp      pop gdpPercap lifeExpMonths
##   <fct>         <fct>    <int> <dbl>    <int>    <dbl>      <dbl>
## 1 Japan          Asia      2007   82.6 1.27e8   31656        991
## 2 Hong Kong, China Asia      2007   82.2 6.98e6   39725        986
## 3 Iceland        Europe      2007   81.8 3.02e5   36181        981
## 4 Switzerland    Europe      2007   81.7 7.55e6   37506        980
## 5 Australia      Oceania     2007   81.2 2.04e7   34435        975
## 6 Spain           Europe      2007   80.9 4.04e7   28821        971
## 7 Sweden          Europe      2007   80.9 9.03e6   33860        971
```

```
## 8 Israel      Asia      2007      80.7 6.43e6      25523      969
## 9 France      Europe     2007      80.7 6.11e7      30470      968
## 10 Canada     Americas  2007      80.7 3.34e7      36319      968
## # ... with 132 more rows
```

#### 4. summarize

summarize returns a tibble with “summarizing” statistics of the dataset you give.

```
# Summarize to find the median life expectancy
```

```
gapminder %>% summarize(medianLifeExp = median(lifeExp))
```

```
## # A tibble: 1 x 1
##   medianLifeExp
##           <dbl>
## 1           60.7
```

```
# Filter for 1957 then summarize the median life expectancy
```

```
gapminder %>% filter(year == 2007) %>%
  summarize(medianLifeExp = median(lifeExp))
```

```
## # A tibble: 1 x 1
##   medianLifeExp
##           <dbl>
## 1           71.9
```

```
# Summarize can even return two values
```

```
gapminder %>% filter(year == 1957) %>%
  summarize(medianLifeExp = median(lifeExp), maxGdpPercap = max(gdpPercap))
```

```
## # A tibble: 1 x 2
##   medianLifeExp maxGdpPercap
##           <dbl>         <dbl>
## 1           48.4       113523
```

#### 5. group\_by

In the previous example, filter allowed you to select a value of one particular column, which then you summarized. In order to summarize over all values of a particular column, we use group\_by. It’s equivalent to using a for loop on filter where the loop runs over the all the possible arguments passable to filter (which is what you would do to get an overall summary table if you didn’t know about group\_by).

```
# Find median life expectancy and maximum GDP per capita in each year
```

```
gapminder %>%
  group_by(year) %>%
  summarize(medianLifeExp = median(lifeExp), maxGdpPercap = max(gdpPercap))
```

```
## # A tibble: 12 x 3
##   year medianLifeExp maxGdpPercap
##   <int>         <dbl>         <dbl>
## 1  1952          45.1         108382
## 2  1957          48.4         113523
## 3  1962          50.9          95458
## 4  1967          53.8          80895
## 5  1972          56.5        109348
## 6  1977          59.7          59265
## 7  1982          62.4          33693
## 8  1987          65.8          31541
## 9  1992          67.7          34933
## 10 1997          69.4          41283
## 11 2002          70.8          44684
## 12 2007          71.9          49357
```

*# Find median life expectancy and maximum GDP per capita in each continent in 1957*

```
gapminder %>%
  filter(year==1957) %>%
  group_by(continent) %>%
  summarize(medianLifeExp = median(lifeExp), maxGdpPercap = max(gdpPercap))
```

```
## # A tibble: 5 x 3
##   continent medianLifeExp maxGdpPercap
##   <fct>         <dbl>         <dbl>
## 1 Africa          40.6          5487
## 2 Americas         56.1        14847
## 3 Asia            48.3        113523
## 4 Europe          67.6        17909
## 5 Oceania         70.3        12247
```

*#You can even pass two arguments to group\_by. In that case it returns the summary table*

```
gapminder %>%
  group_by(continent, year) %>%
  summarize(medianLifeExp = median(lifeExp), maxGdpPercap = max(gdpPercap))
```

```
## # A tibble: 60 x 4
## # Groups:   continent [?]
##   continent  year medianLifeExp maxGdpPercap
##   <fct>    <int>         <dbl>         <dbl>
## 1 Africa   1952          38.8          4725
## 2 Africa   1957          40.6          5487
## 3 Africa   1962          42.6          6757
## 4 Africa   1967          44.7         18773
```

```
## 5 Africa      1972      47.0      21011
## 6 Africa      1977      49.3      21951
## 7 Africa      1982      50.8      17364
## 8 Africa      1987      51.6      11864
## 9 Africa      1992      52.4      13522
## 10 Africa     1997      52.8      14723
## # ... with 50 more rows
```

*# In the above output, summarize gave the summaries at all possible (year, continent)*

## Visualization with ggplot2

ggplot2 is the canonical data visualization tool. We'll load it here


```
library("ggplot2")
```

The basic syntax for making a plot is `ggplot(#dataframe, aes(x=#xcomponent, y=#component))`. The `+` operator will signify additional features to the plot.


### 1. Scatter plot

Scatter plots are achieved using `+ geom_point()`.

```
library(dplyr)
library(gapminder)
gapminder_1952 <- gapminder %>%
  filter(year == 1952)
ggplot(gapminder_1952, aes(x = pop, y = gdpPercap)) +
  geom_point()
```




```
ggplot(gapminder_1952, aes(x=pop, y=lifeExp)) + geom_point()
```



## 2. Log scale


You can use log axes by adding `+ scale_x_log10()` or `+ scale_y_log10()`.

```
ggplot(gapminder_1952, aes(x = pop, y = lifeExp)) +  
  geom_point() +  
  scale_x_log10()
```



```
ggplot(gapminder_1952, aes(x=pop, y=gdpPercap)) +  
  geom_point() +  
  scale_x_log10() +  
  scale_y_log10()
```







### 3. Aesthetics

You can add more than just `x` and `y` in the `aes()` parameter in `ggplot`. For example, you can add `color` which is the colour of the data points, and `size` which is their size. For colour, `ggplot` automatically adds a legend!

```
ggplot(gapminder_1952, aes(x=pop, y=lifeExp, color=continent)) +  
  geom_point() +  
  scale_x_log10()
```



```
ggplot(gapminder_1952, aes(x = pop, y = lifeExp, color = continent, size = gdpPercap)) +  
  geom_point() +  
  scale_x_log10()
```



Question: do **color** and **size** only matter when you use the scatter plot? Question: is it possible to modify the legend (say represent the **gdpPercap** values in log scale?)


#### 4. Faceting

Faceting refers to the practice of making multiple plots within the same figure. In **ggplot** you can make multiple plots according to their indices in the dataset. To facet a plot, we use `+ facet_wrap(~#indexingparameter)`.

*# Here we break down the analysis of population vs life expectancy according to continent*

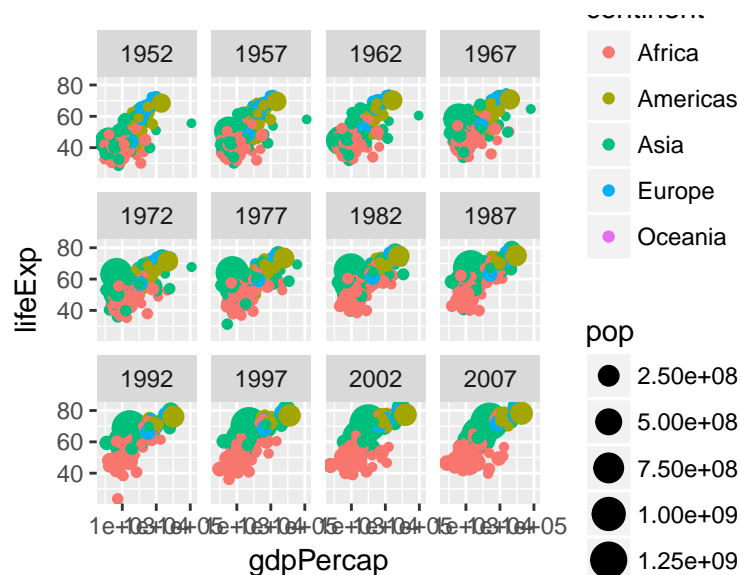
*# Scatter plot comparing pop and lifeExp, faceted by continent*

```
ggplot(gapminder_1952, aes(x=pop, y=lifeExp, size = pop, color = continent)) +  
  geom_point() +  
  scale_x_log10() +  
  facet_wrap(~ continent)
```



Note that we can apply the `color` and the `facet_wrap` functions to the same parameters. The result is obviously that all the colours go into the same subplot.

```
# Scatter plot comparing gdpPercap and lifeExp, with color representing continent
# and size representing population, faceted by year
ggplot(gapminder, aes(x=gdpPercap, y=lifeExp, color = continent, size = pop)) +
  geom_point() +
  scale_x_log10() +
  facet_wrap(~ year)
```



## 5. Axes limits

Set axes limits using `+ expand_limits(y=#value)`


```

# Here we'll visualize some summarized data
by_year <- gapminder %>%
  group_by(year) %>%
  summarize(medianLifeExp = median(lifeExp),
            maxGdpPerCap = max(gdpPerCap))

# Create a scatter plot showing the change in medianLifeExp over time

ggplot(by_year, aes(x = year, y = medianLifeExp)) +
  geom_point() +
  expand_limits(y=0)

```




```

# Summarize medianGdpPerCap within each continent within each year: by_year_continent
by_year_continent <- gapminder %>%
  group_by(year, continent) %>%
  summarize(medianGdpPerCap = median(gdpPerCap))

# Plot the change in medianGdpPerCap in each continent over time

ggplot(by_year_continent, aes(x=year, y=medianGdpPerCap, color = continent)) +
  geom_point() +
  expand_limits(y=0)

```




## 6. Line plot

To make a Line plot, use + `geom_line()`

```
# Summarize the median gdpPercap by year, then save it as by_year
by_year <- gapminder %>%
  group_by(year) %>%
  summarize(medianGdpPercap = median(gdpPercap))

# Create a line plot showing the change in medianGdpPercap over time

ggplot(by_year, aes(x=year, y=medianGdpPercap)) +
  geom_line() +
  expand_limits(y=0)
```




```

# Summarize the median gdpPercap by year & continent, save as by_year_continent
by_year_continent <- gapminder %>%
  group_by(year, continent) %>%
  summarize(medianGdpPercap = median(gdpPercap))

# Create a line plot showing the change in medianGdpPercap by continent over time

ggplot(by_year_continent, aes(x=year, y=medianGdpPercap, color = continent)) +
  geom_line() +
  expand_limits(y=0)

```



## 7. Bar plot

Use + `geom_col()`


```

# Summarize the median gdpPercap by year and continent in 1952
by_continent <- gapminder %>%
  filter(year==1952) %>%
  group_by(continent) %>%
  summarize(medianGdpPercap = median(gdpPercap))

# Create a bar plot showing medianGdp by continent

ggplot(by_continent, aes(x=continent, y=medianGdpPercap)) +
  geom_col()

```



## 8. Histogram


Use + `geom_histogram(binwidth=#setwidth)`. Note that for hist, you need only specify `x` in `aes()`.

```
gapminder_1952 <- gapminder %>%
  filter(year == 1952)

# Create a histogram of population (pop), with x on a log scale

ggplot(gapminder_1952, aes(x=pop)) +
  scale_x_log10() +
  geom_histogram()
```


## ``stat_bin()`` using ``bins = 30``. Pick better value with ``binwidth``.



## 9. Box-Whisker plot

Use + `geom_boxplot()`

```
gapminder_1952 <- gapminder %>%  
  filter(year == 1952)  
  
# Create a boxplot comparing gdpPercap among continents  
  
ggplot(gapminder_1952, aes(x=continent, y=gdpPercap)) +  
  geom_boxplot() +  
  scale_y_log10()
```



## 10. Adding title and x and y labels

For title use + `ggtitle()` and for labels use `labs(x=#namehere, y=#namehere)`

```
gapminder_1952 <- gapminder %>%  
  filter(year == 1952)  
  
# Add a title to this graph: "Comparing GDP per capita across continents"  
ggplot(gapminder_1952, aes(x = continent, y = gdpPercap)) +  
  geom_boxplot() +  
  scale_y_log10() +  
  ggtitle("Comparing GDP per capita across continents") +  
  labs(x="Continent", y="GDP per Capita")
```



Comparing GDP per capita across contin

