

Bayesian linear regression

Jingchen (Monika) Hu

MATH 347 Bayesian Statistics

Installing the necessary packages

```
install.packages("devtools")
require(devtools)
devtools::install_github("bayesball/ProbBayes")
require(utils)
require(ggplot2)
require(gridExtra)
require(ProbBayes)
require(tidyverse)
crcblue <- "#2905a1"
```


Introduction: Adding a continuous predictor variable

The simple linear regression model

```
library(readr)
CEData <- read_csv("CEsample.csv")

## Parsed with column specification:
## cols(
##   UrbanRural = col_double(),
##   TotalIncomeLastYear = col_double(),
##   Race = col_double(),
##   TotalExpLastQ = col_double(),
##   log_TotalIncome = col_double(),
##   log_TotalExp = col_double(),
##   Rural = col_double()
## )

g1 <- ggplot(CEData, aes(x = log_TotalIncome, y = log_TotalExp)) +
  geom_point(size=1) +
  labs(x = "log(Income)", y = "log(Expenditure)") +
  theme_grey(base_size = 10, base_family = "")
g1
```



The CE sample

A simple linear regression for the CE sample

MCMC simulation by JAGS for the SLR model

JAGS script for the SLR model

```
modelString <-"
model {
  ## sampling
  for (i in 1:N){
    y[i] ~ dnorm(beta0 + beta1*x[i], invsigma2)
  }

  ## priors
  beta0 ~ dnorm(mu0, g0)
  beta1 ~ dnorm(mu1, g1)
  invsigma2 ~ dgamma(a, b)
  sigma <- sqrt(pow(invsigma2, -1))
}
```

```
}
"
```

- Pass the data and hyperparameter values to JAGS:

```
y <- as.vector(CEDData$log_TotalExp)
x <- as.vector(CEDData$log_TotalIncome)
N <- length(y)
the_data <- list("y" = y, "x" = x, "N" = N,
                "mu0" = 0, "g0" = 0.0001,
                "mu1" = 0, "g1" = 0.0001,
                "a" = 1, "b" = 1)

initsfunction <- function(chain){
  .RNG.seed <- c(1,2)[chain]
  .RNG.name <- c("base::Super-Duper",
                "base::Wichmann-Hill")[chain]
  return(list(.RNG.seed=.RNG.seed,
              .RNG.name=.RNG.name))
}
```

- Run the JAGS code for this model:

```
posterior <- run.jags(modelString,
                     n.chains = 1,
                     data = the_data,
                     monitor = c("beta0", "beta1", "sigma"),
                     adapt = 1000,
                     burnin = 5000,
                     sample = 5000,
                     thin = 1,
                     inits = initsfunction)
```

```
## Calling the simulation...
## Welcome to JAGS 4.3.0 on Sun Nov 24 21:04:09 2019
## JAGS is free software and comes with ABSOLUTELY NO WARRANTY
## Loading module: basemod: ok
## Loading module: bugs: ok
## . . Reading data file data.txt
## . Compiling model graph
##   Resolving undeclared variables
##   Allocating nodes
## Graph information:
##   Observed stochastic nodes: 994
##   Unobserved stochastic nodes: 3
##   Total graph size: 3466
## . Reading parameter file inits1.txt
## . Initializing model
## . Adaptation skipped: model is not in adaptive mode.
```

```
## . Updating 5000
## -----| 5000
## ***** 100%
## . . . . Updating 5000
## -----| 5000
## ***** 100%
## . . . . Updating 0
## . Deleting model
## .
## Note: the model did not require adaptation
## Simulation complete. Reading coda files...
## Coda files loaded successfully
## Calculating summary statistics...
## Finished running the simulation
```

JAGS output for the SLR model





- Obtain posterior summaries of all parameters:

```
summary(posterior)
```

```
##      Lower95   Median Upper95      Mean      SD Mode      MCerr
## beta0 3.771420 4.3136450 4.655760 4.2947731 0.22275382   NA 0.0421194486
## beta1 0.388552 0.4218125 0.471380 0.4237176 0.02091679   NA 0.0039634093
## sigma 0.694060 0.7249680 0.757081 0.7251875 0.01624723   NA 0.0002175128
##      MC%ofSD SSeff      AC.10 psrf
## beta0    18.9    28 0.89361268   NA
## beta1    18.9    28 0.89351502   NA
## sigma     1.3  5579 0.02058294   NA
```





```
plot(posterior, vars = "beta0")
```

```
## Generating plots...
```



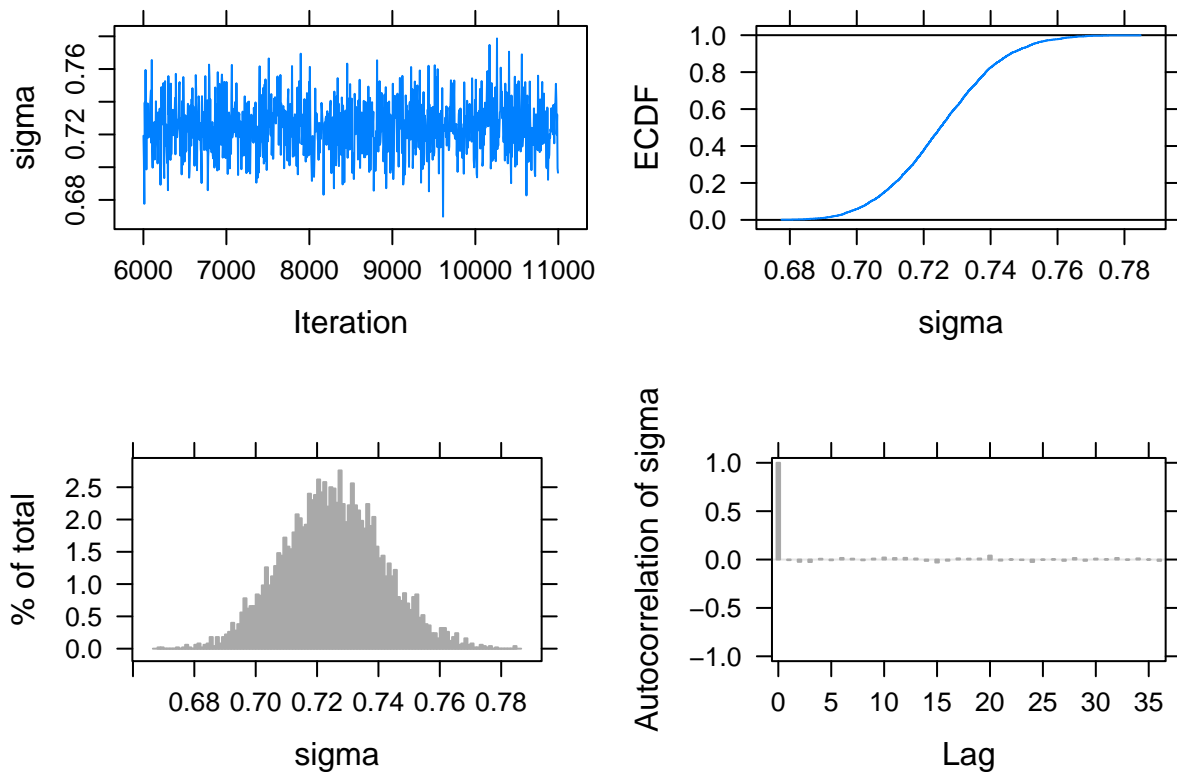
```
plot(posterior, vars = "beta1")
```

Generating plots...



```
plot(posterior, vars = "sigma")
```

```
## Generating plots...
```



New JAGS script for the SLR model

Setting `thin = 50`, to get rid of the stickiness in β_0 and β_1 .

```
posterior_new <- run.jags(modelString,
  n.chains = 1,
  data = the_data,
  monitor = c("beta0", "beta1", "sigma"),
  adapt = 1000,
  burnin = 5000,
  sample = 5000,
  thin = 50,
  inits = initsfunction)
```

```
## Calling the simulation...
## Welcome to JAGS 4.3.0 on Sun Nov 24 21:04:15 2019
## JAGS is free software and comes with ABSOLUTELY NO WARRANTY
## Loading module: basemod: ok
## Loading module: bugs: ok
## . . Reading data file data.txt
## . Compiling model graph
##   Resolving undeclared variables
```

```

##      Allocating nodes
## Graph information:
##      Observed stochastic nodes: 994
##      Unobserved stochastic nodes: 3
##      Total graph size: 3466
## . Reading parameter file inits1.txt
## . Initializing model
## . Adaptation skipped: model is not in adaptive mode.
## . Updating 5000
## -----| 5000
## ***** 100%
## . . . . Updating 250000
## -----| 250000
## ***** 100%
## . . . . Updating 0
## . Deleting model
## .
## Note: the model did not require adaptation
## Simulation complete. Reading coda files...
## Coda files loaded successfully
## Calculating summary statistics...
## Finished running the simulation

```

New JAGS output for the SLR model

- Obtain posterior summaries of all parameters:

```
summary(posterior_new)
```





```

##      Lower95   Median Upper95      Mean      SD Mode      MCerr
## beta0 3.926200 4.3251800 4.753090 4.3269816 0.21092686  NA 0.0054707075
## beta1 0.381057 0.4208825 0.458968 0.4207384 0.01977946  NA 0.0004930936
## sigma 0.693623 0.7254665 0.757397 0.7255765 0.01624891  NA 0.0002297942
##      MC%ofSD SSeff      AC.500 psrf
## beta0      2.6  1487 0.018912301  NA
## beta1      2.5  1609 0.018555972  NA
## sigma      1.4  5000 0.003076245  NA

```





```
plot(posterior_new, vars = "beta0")
```

```
## Generating plots...
```



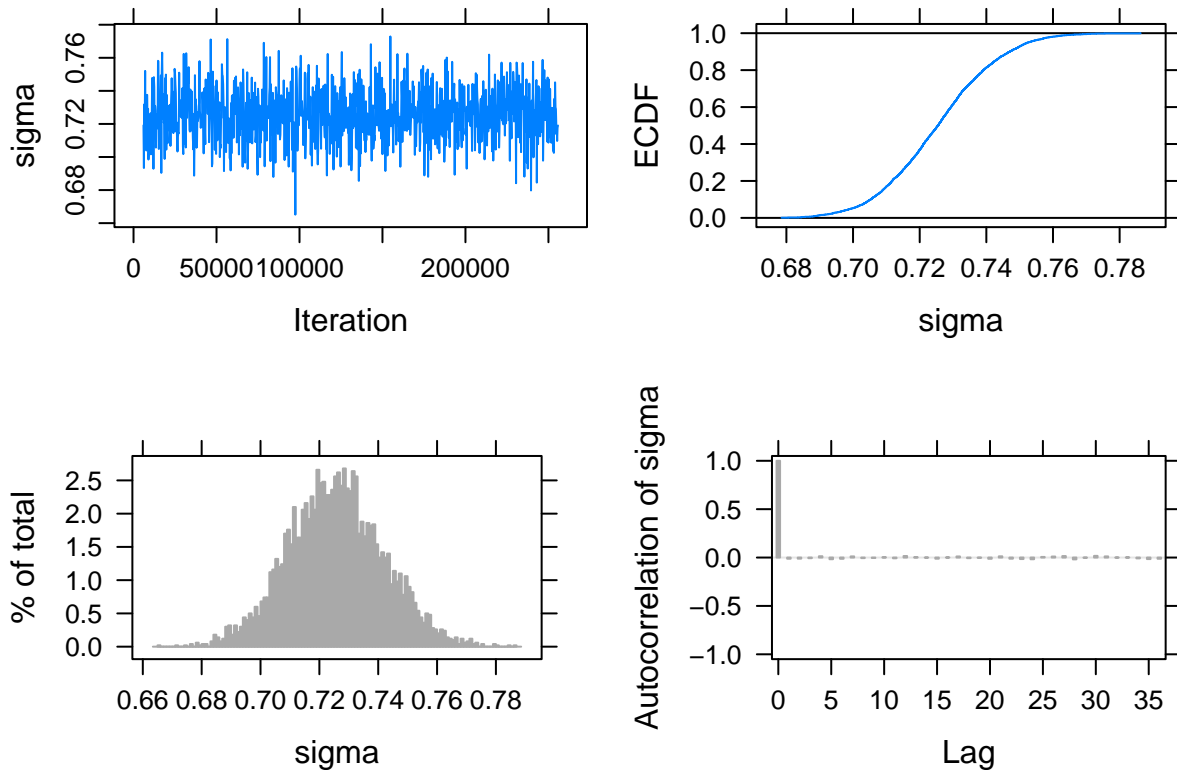
```
plot(posterior_new, vars = "beta1")
```

Generating plots...




```
plot(posterior_new, vars = "sigma")
```

```
## Generating plots...
```




Bayesian inferences with SLR

Simulate fits from the regression model

```
post <- as.mcmc(posterior_new)
post_means <- apply(post, 2, mean)
post <- as.data.frame(post)
```

```
ggplot(CEDData, aes(log_TotalIncome, log_TotalExp)) +
  geom_point(size=1) +
  geom_abline(data=post[1:10, ],
             aes(intercept=beta0, slope=beta1), alpha = 0.5) +
  geom_abline(intercept = post_means[1],
             slope = post_means[2], size = 1) +
  ylab("log(Expenditure)") + xlab("log(Income)") +
  theme_grey(base_size = 10, base_family = "")
```




Learning about the expected response


```
post <- as.data.frame(post)
one_expected <- function(x){
  lp <- post[, "beta0"] + x * post[, "beta1"]
  data.frame(Value = paste("log(Income) =", x),
             Expected_logExp = lp)
}

df <- map_df(c(1, 5, 7, 9), one_expected)

require(ggribes)
ggplot(df, aes(x = Expected_logExp, y = Value)) +
  geom_density_ridges() +
  theme_grey(base_size = 8, base_family = "")
```



```
df <- map_df(c(10, 11, 12, 13), one_expected)
ggplot(df, aes(x = Expected_logExp, y = Value)) +
  geom_density_ridges() +
  theme_grey(base_size = 8, base_family = "")
```




```
df %>% group_by(Value) %>%
  summarize(P05 = quantile(Expected_logExp, 0.05),
            P50 = median(Expected_logExp),
            P95 = quantile(Expected_logExp, 0.95))
```

```
## # A tibble: 4 x 4
##   Value          P05    P50    P95
##   <chr>        <dbl> <dbl> <dbl>
## 1 log(Income) = 10  8.49  8.53  8.58
## 2 log(Income) = 11  8.91  8.96  8.99
## 3 log(Income) = 12  9.32  9.38  9.43
## 4 log(Income) = 13  9.71  9.80  9.88
```

Prediction of future responses

```
one_predicted <- function(x){
  lp <- post[, "beta0"] + x * post[, "beta1"]
  y <- rnorm(5000, lp, post[, "sigma"])
  data.frame(Value = paste("Size =", x),
             Predicted_logExp = y)
}
df <- map_df(c(1, 5, 7, 9), one_predicted)
```

```
require(ggribes)
ggplot(df, aes(x = Predicted_logExp, y = Value)) +
  geom_density_ridges() +
  theme_grey(base_size = 9, base_family = "")
```



```
df %>% group_by(Value) %>%
  summarize(P05 = quantile(Predicted_logExp, 0.05),
            P50 = median(Predicted_logExp),
            P95 = quantile(Predicted_logExp, 0.95))
```


```
## # A tibble: 4 x 4
##   Value      P05    P50    P95
##   <chr>    <dbl> <dbl> <dbl>
## 1 Size = 1  3.52  4.75  5.95
## 2 Size = 5  5.23  6.44  7.61
## 3 Size = 7  6.06  7.25  8.51
## 4 Size = 9  6.90  8.12  9.32
```

More on priors

Subjective prior: standardization

```
CEData$log_TotalExpSTD <- scale(CEData$log_TotalExp)
CEData$log_TotalIncomeSTD <- scale(CEData$log_TotalIncome)
```

```
g2 = ggplot(CEData, aes(x = log_TotalIncomeSTD, y = log_TotalExpSTD)) +
  geom_point(size=1) +
  xlab("log(Income) STD") + ylab("log(Expenditure) STD") +
  theme_grey(base_size = 10, base_family = "")
grid.arrange(g1, g2, ncol=2)
```



Subjective prior: JAGS script for the standardized SLR model

```
modelString <- "
model {
  ## sampling
  for (i in 1:N){
    y[i] ~ dnorm(beta0 + beta1*x[i], invsigma2)
  }
}
```

```
## priors
beta0 ~ dnorm(mu0, g0)
beta1 ~ dnorm(mu1, g1)
invsigma2 ~ dgamma(a, b)
sigma <- sqrt(pow(invsigma2, -1))
}
"
```

- Pass the data and hyperparameter values to JAGS:

```
y <- as.vector(CEDData$log_TotalExpSTD)
x <- as.vector(CEDData$log_TotalIncomeSTD)
N <- length(y)
the_data <- list("y" = y, "x" = x, "N" = N,
                 "mu0" = 0, "g0" = 1,
                 "mu1" = 0.7, "g1" = 44.4,
                 "a" = 1, "b" = 1)

initsfunction <- function(chain){
  .RNG.seed <- c(1,2)[chain]
  .RNG.name <- c("base::Super-Duper",
                "base::Wichmann-Hill")[chain]
  return(list(.RNG.seed=.RNG.seed,
              .RNG.name=.RNG.name))
}
```

- Run the JAGS code for this model:

```
posterior_sub <- run.jags(modelString,
                          n.chains = 1,
                          data = the_data,
                          monitor = c("beta0", "beta1", "sigma"),
                          adapt = 1000,
                          burnin = 5000,
                          sample = 5000,
                          thin = 1,
                          inits = initsfunction)
```

```
## Calling the simulation...
## Welcome to JAGS 4.3.0 on Sun Nov 24 21:04:58 2019
## JAGS is free software and comes with ABSOLUTELY NO WARRANTY
## Loading module: basemod: ok
## Loading module: bugs: ok
## . . Reading data file data.txt
## . Compiling model graph
##   Resolving undeclared variables
##   Allocating nodes
## Graph information:
##   Observed stochastic nodes: 994
```

```

##      Unobserved stochastic nodes: 3
##      Total graph size: 3466
## . Reading parameter file inits1.txt
## . Initializing model
## . Adaptation skipped: model is not in adaptive mode.
## . Updating 5000
## -----| 5000
## ***** 100%
## . . . . Updating 5000
## -----| 5000
## ***** 100%
## . . . . Updating 0
## . Deleting model
## .
## Note: the model did not require adaptation
## Simulation complete. Reading coda files...
## Coda files loaded successfully
## Calculating summary statistics...
## Finished running the simulation

```

Subjective prior: JAGS output for the SLR model

- Obtain posterior summaries of all parameters:

```
summary(posterior_sub)
```

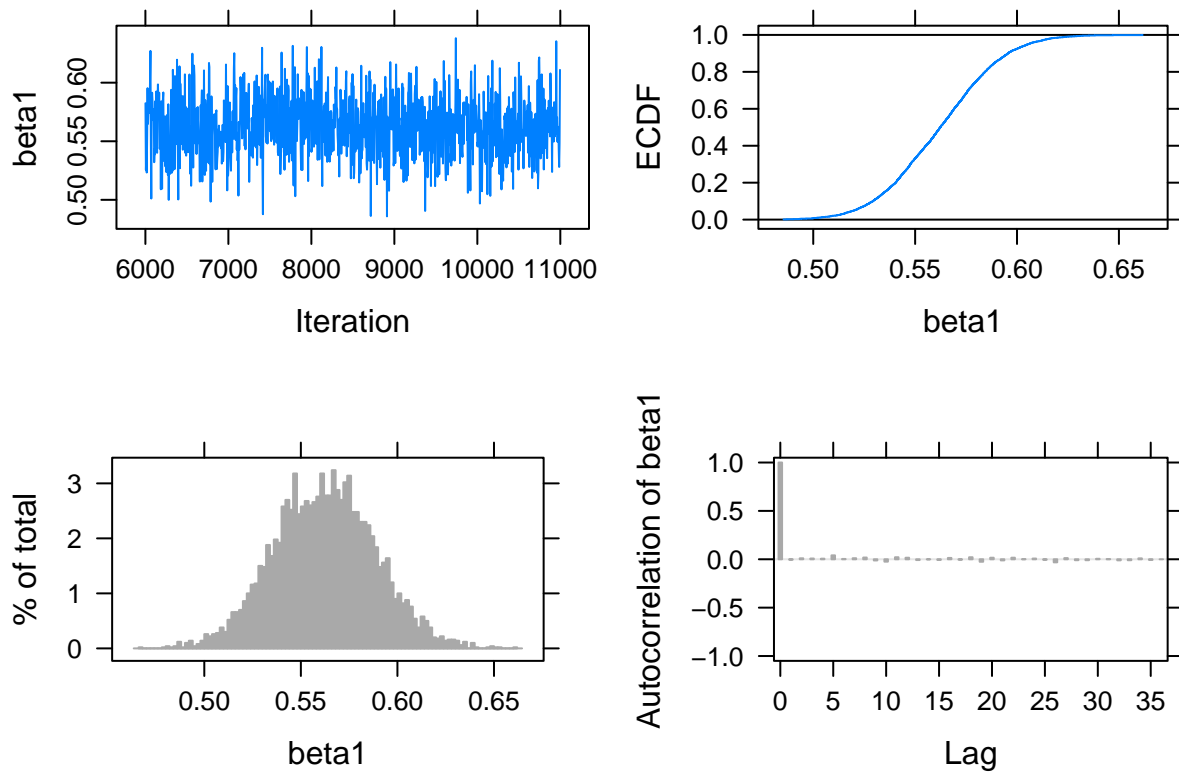
```

##           Lower95      Median  Upper95      Mean      SD Mode
## beta0 -0.0484962 -0.0000908726 0.0544673 -5.334241e-05 0.02637737 NA
## beta1  0.5131890  0.5623200000 0.6144910  5.622705e-01 0.02616385 NA
## sigma  0.7958380  0.8314665000 0.8684720  8.318869e-01 0.01866642 NA
##
##           MCerr MC%ofSD SSeff      AC.10 psrf
## beta0 0.0003561888      1.4  5484 -0.02064139  NA
## beta1 0.0003700127      1.4  5000 -0.02451218  NA
## sigma 0.0002502895      1.3  5562  0.02083092  NA

```

```
plot(posterior_sub, vars = "beta1")
```

```
## Generating plots...
```

Conditional means prior: JAGS script

```
modelString <- "
model {
  ## sampling
  for (i in 1:N){
    y[i] ~ dnorm(beta0 + beta1*x[i], invsigma2)
  }

  ## priors
  beta1 <- (mu2 - mu1)/(x2 - x1)
  beta0 <- mu1 - x1*(mu2 - mu1)/(x2 - x1)
  mu1 ~ dnorm(m1, g1)
  mu2 ~ dnorm(m2, g2)
  invsigma2 ~ dgamma(a, b)
  sigma <- sqrt(pow(invsigma2, -1))
}
"
```